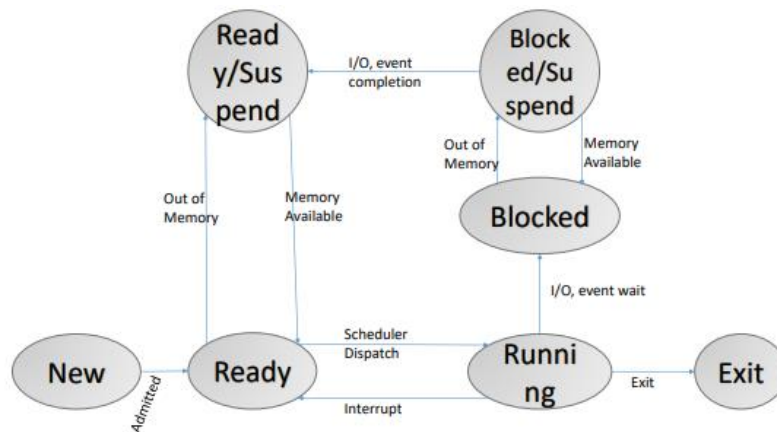


Scheduling

- cpu burst is followed by an io burst
- cpu scheduling decisions may take place when a process:
 - switches from running to waiting state
 - switches from running to ready state
 - wait to ready
 - when a process exits



- cpu scheduler : the program which selects which process in the ready queue to load to CPU
- Dispatcher : Program which loads the selected program to the CPU
- Dispatch latency : time between stopping one process and loading the next.

Pre-emption

- If the current running process can be removed by another process it's pre-emptive
 - If it can be removed forcefully, it is non-cooperative
- Race condition can occur in a preemptive system.
 - When data is shared across several processes.

Dispatcher

- Switching context
- Switching to user mode
- Jumping to the proper location in the user program to restart that program.

Scheduling criteria



Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible (40% - 100%) real system
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process (time to get info mem + time in queues + time execution + time I/O) *final output*
- **Waiting time** – amount of time a process has been waiting in the ready queue *As algorithm only affects the waiting time*
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced. *not output - just response*
 (better alternative to turnaround time to use in interactive systems)
 as TA time is dependent on the I/O speed
 TA needs full completion
 Response does not need that



Scheduling algorithms

- FCFS - [Picture on page "Chapter 05 - CPU Scheduling"](#) ([Web view](#))
 - If two processors arrive at the same time, the one with smaller pid gets the cpu first
 - Convoy – Short processes behind long processes
 - Not good for time sharing systems
- Shortest job first (sjf) - [Picture on page "Chapter 05 - CPU Scheduling"](#) ([Web view](#))
 - Can only estimate the length of the burst time
- Shortest remaining job first
 - Preemptive version of the above algorithm
 - Behave same as sjf when all the process are arrived to the queue
- Round robin
 - Time sharing system
 - A time quantum is assigned to each thread
 - if q is large -> FCFS
 - if q is small -> overhead
- priority scheduling
 - equal priority -> FCFS
 - every algorithm is a priority scheduling algorithm
 - starvation can occur – solution : aging
 - multilevel queue
 - prioritization based upon process type
 - multilevel feedback queue
 - process can move between multiple queues
 - queues can have different scheduling algorithms
 - aging is implemented in this manner

- typically done via priority set by programmer



Multiple-Processor Scheduling – Load Balancing

- If SMP, need to keep all CPUs loaded for efficiency
- **Load balancing** attempts to keep workload evenly distributed
- **Push migration** – periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs
- **Pull migration** – idle processors pulls waiting task from busy processor



Real-time scheduling

- soft real time – no enforcement / no deadline
- hard real-time – task must be serviced by its deadline
- latency
 - interrupt latency – time from arrival of interrupt to start of routine that service interrupt
 - dispatch latency – time for schedule to take current process off CPU and switch to another
-
-
-
-
-
-

Operating Systems

The major differences between a process and a thread are given as follows –

| Comparison Basis | Process | Thread |
|------------------------|---|---|
| Definition | A process is a program under execution i.e an active program. | A thread is a lightweight process that can be managed independently by a scheduler. |
| Context switching time | Processes require more time for context switching as they are more heavy. | Threads require less time for context switching as they are lighter than processes. |
| Memory Sharing | Processes are totally independent and don't share memory. | A thread may share some memory with its peer threads. |
| Communication | Communication between processes requires more time than between threads. | Communication between threads requires less time than between processes . |
| Blocked | If a process gets blocked, remaining processes can continue execution. | If a user level thread gets blocked, all of its peer threads also get blocked. |
| Resource Consumption | Processes require more resources than threads. | Threads generally need less resources than processes. |
| Dependency | Individual processes are independent of each other. | Threads are parts of a process and so are dependent. |
| Data and Code sharing | Processes have independent data and code segments. | A thread shares the data segment, code segment, files etc. with its peer threads. |
| Treatment by OS | All the different processes are treated separately by the operating system. | All user level peer threads are treated as a single task by the operating system. |
| Time for creation | Processes require more time for creation. | Threads require less time for creation. |
| Time for termination | Processes require more time for termination. | Threads require less time for termination. |

File System

- Directory Structure
 - Nodes containing information about all files.
- File operation
 - Create
 - Write
 - Read
 - Reposition within file – seek
 - Delete
 - Truncate
 - Open
 - Close
- Access methods
 - Sequential – simplest
 - Going step by step
 - Read next, write next and reset
 - Direct
 - Can go to the specific location in the memory
 - read n, write n,

Operating Systems

- position to n
 - read next
 - write next
 - rewrite n; n – relative block number
- Sequential access is also possible
- Other access methods

Disk structure

- Two types
 - General file-system
 - Special purpose file system




Operations Performed on Directory

- Search for a file
- Create a file
- Delete a file
- List a directory
- Rename a file
- Traverse the file system

- Directory organization goals
 - Efficiency
 - Naming
 - Grouping
- Types of directories
 - Single level – single directory for all users
 - Limitations when files and users increase
 - Naming problem/ grouping problem
 - Two level
 - Separate directory for each user
 - Path name available
 - Search efficiency
 - No grouping
 - Tree structured
 - Most common
 - Path name available for files
 - Acyclic graph
 - Shared subdirectories and files are available
 - One copy, accessed by all users
 - Limitations
 - Dangling point when one user deletes a shared file and the other is not updated about the file
 - Solution – back pointers using a daisy chain

Daisy Chain

Electrical Engineering



In electrical and electronic engineering, a daisy chain is a wiring scheme in which multiple devices are wired together in sequence or in a ring, similar to a garland of daisy flowers. Daisy chains may be used for power, analog signals, digital data, or a combination of them.

-
- Link – a pointer to an existing file – ignores traversing through the tree.
- Resolve the link - path
- General graph directory
 - Possible defects
 - Self-referencing
 - Cycles –
- How do we guarantee no cycles?
 - Allow only links to files not subdirectories
 - **Garbage collection**
 - Every time a new link is added use a cycle detection algorithm to determine whether it is OK

very time consuming for a large directory

- File-system Structure
 - Resides on secondary storage
 - File control block – Storage structure consisting of information about a file
- File System Layers
 - Devices
 - I/O control – Device drivers manage devices
 - Basic File system –
 - Translates commands to I/O control
 - Manages memory buffers (Hold data in transit) and caches (Hold frequently used)
 - File organization module –
 - translates logical block to physical block
 - Manages free space, disk allocation
 - Logical file system –
 - Manages metadata information
 - Directory management
 - Protection
 - Application programs
- Boot control block – contains info needed to boot OS from a volume
- Volume control block – contains information about the volume
- File control block –
 - One FCB per file
 - Contains information about the file
- In-memory file system structures
 - Mount tables – Store file system mounts, mount points, file system types
 - System-wide open-file table – contains a copy of FCB of open files
 - Per-process open-file table – contains a pointer to appropriate entries in System wide table
- Directory Implementation

Operating Systems

- Linear list
 - Simple to program
 - Not efficient
- Hash table
 - Decrease search time
 - Collisions – two names can hash to same location
 - Good only if entries are fixed sized
 - Else use chained overflow
- Memory allocation method
 - Contiguous
 - Best performance in most cases
 - Simple
 - Problems
 - Finding space in disk
 - Knowing file size
 - External fragmentation
 - Extent based systems – modified contiguous
 - Allocate disk blocks in extents
 - Extent is a contiguous block of disks
 - A file consists of one or more extents
 - Linked
 - Each file is a linked list of blocks.
 - File ends at nil pointer
 - No external fragmentation
 - Reliability is less
 - File Allocation Table
 - Indexed allocation method
 - Each file has an index table with pointers to its data blocks.
- Performance
 - Contiguous – good for sequential and random
 - Linked – good for sequential
 - Indexed – more complex
 - Clustering can improve performance
- Free-space management
 - Maintain a free-space list
 - A bit vector.
 - Easy to get contiguous files
 - Linked list (Free list)
 - No waste
 - Grouping
 - Counting
 - Keep address of first free block and next number of free blocks
 - Free space list has entries containing addresses and counts
 - Space maps
 - Used in ZFS
 - Divides device space into **metaslab** units
 - Each metaslab has associated space map
 - Records log to a file instead of filesystem
 - Combines contiguous free blocks into single entry
- TRIMming unused blocks
 - Never mechanism to inform NVM devices that a page is free

Operating Systems

- Then can be garbage collected or erased
 - Efficiency and Performance
 - Efficiency Depends on
 - Disk allocation
 - Directory algorithms
 - Disk allocation and directory algorithms
 - Types of data kept in file's directory entry
 - Pre-allocation or as-needed allocation of metadata structures
 - Fixed-size or varying-size data structures
-
- Performance
 - Keeping data and metadata close
 - Page cache
 - Caches pages rather than disk blocks
 - Memory mapped I/O uses page cache
 - Routine I/O uses buffer (Disk) cache
 - Unified buffer cache
 - Use the same page cache in Memory-mapped I/O and Ordinary I/O
 - Avoid double caching
 - Recovery
 - Consistency checking compares directory structure with data blocks
 - Slow
 - Can fail
 - Log structured file systems
 - Record each metadata update as a transaction
 - All transactions are written to a log
 - Faster recovery from crashes

File system Internals

- Partitions and mounting
 - Root partition – contains the OS, mounted at boot time
 - Other - can be mounted automatically or manually on mount points
 - File system consistency checked at mount time
- Virtual File systems
 - Object oriented way of file systems
 - Can use the same API for different file systems
 - Linux has 4 object types in VFS
 - Inode
 - File
 - Superblock
 - Dentry
- Remote file systems
 - Sharing files across a network
 - Method 1 – FTP (Manually sharing each file)

Operating Systems

- Method 2 – Distributed File System (DFS)
 - Remote directories are visible from the local machine
- Method 3 – WWW
 - Use browsers to locate files
 - Has anonymous access
- Client – server model
- Distributed Information systems (DNS, NIS, CIFS)

- Consistency semantics
 - For evaluating file-sharing systems
 - File sessions

Synchronization Tools

Background

- Concurrent access may result in data inconsistency
- Processes can either
 - Share a logical address space (Code + data)
 - Be allowed to share data through only files or messages

Critical Section Problem

- Changing common data
- When one process is in critical section, no other process should be in critical section
- Permission requested at entry section
- Requirements
 - Mutual Exclusion – No two process
 - Progress – No process should be starved
 - Bounded waiting – When one process requests access to enter critical section, there should be a limit that other processes are allowed to enter their critical section before this one

Peterson's Solution

- Two process solution
- Assume load and store are atomic
- Two processes share two variables turn, flag[2]
- Humble algorithm
- All three requirements are met
- Drawbacks
 - Cannot be used for more than two processes.
 - Cannot be used in multithreaded operations.
- Use a memory barrier to optimize Peterson's
 - Memory models – Guarantees that computer architecture provides to application programs,
 - Strongly ordered – Memory modification of one processor is immediately visible to all other processes.
 - Weakly ordered – may not be immediately visible.
 - Memory barrier – Enforces those changes should be visible to other processes.

Hardware support for synchronization

- Uniprocessors – Could disable interrupts
- Others – Three forms of support
 - Hardware instructions
 - Test and set – Does not solve critical section
 - Returns the original value of the parameter
 - Set the new value to true
 - No bounded wait
 - Can cause starvations
 - Compare and swap
 - Returns the original value of the parameter
 - Set new value if the current is expected
 - Atomic variables
 - Gives atomic updates on basic data types

Mutex Locks

- Simplest

Operating Systems

- Boolean variable – locked or not
- Requires busy waiting
- Acquire and release locks using hardware atomic instructions

Semaphores

- Integer variable
- Wait and signal operations
- Types
 - Counting – To use over a resource having multiple instances. Initialized to the number of instances.
 - Binary – same as mutex
- Implementation with no busy wait
 - Have a waiting queue with each semaphore.
 - Block and wait operations
 - Can have deadlocks and starvations.

Monitors

- A higher-level abstraction. (Itself is an abstract data type)
- Only one process may be active within the monitor at a time. (Mutual exclusion within the monitor)
- Shared variables are defined within the monitor
- Procedures within a monitor can access only local variables and vice versa.

Liveness

Evaluation